

Right-Sizing and Auto-Scaling of MySQL Containers in Kubernetes

Yuan Chen, Min Li

JD.com



About JD.com

China's largest online and overall retailer and biggest Internet company by revenue

- 300 million+ active users
- 2018 revenue: \$67.2 billion

China's largest e-commerce logistics infrastructure and unrivalled nationwide fulfillment network

- 550+ warehouses
- Covering 99% of population
- Standard same-and next day delivery


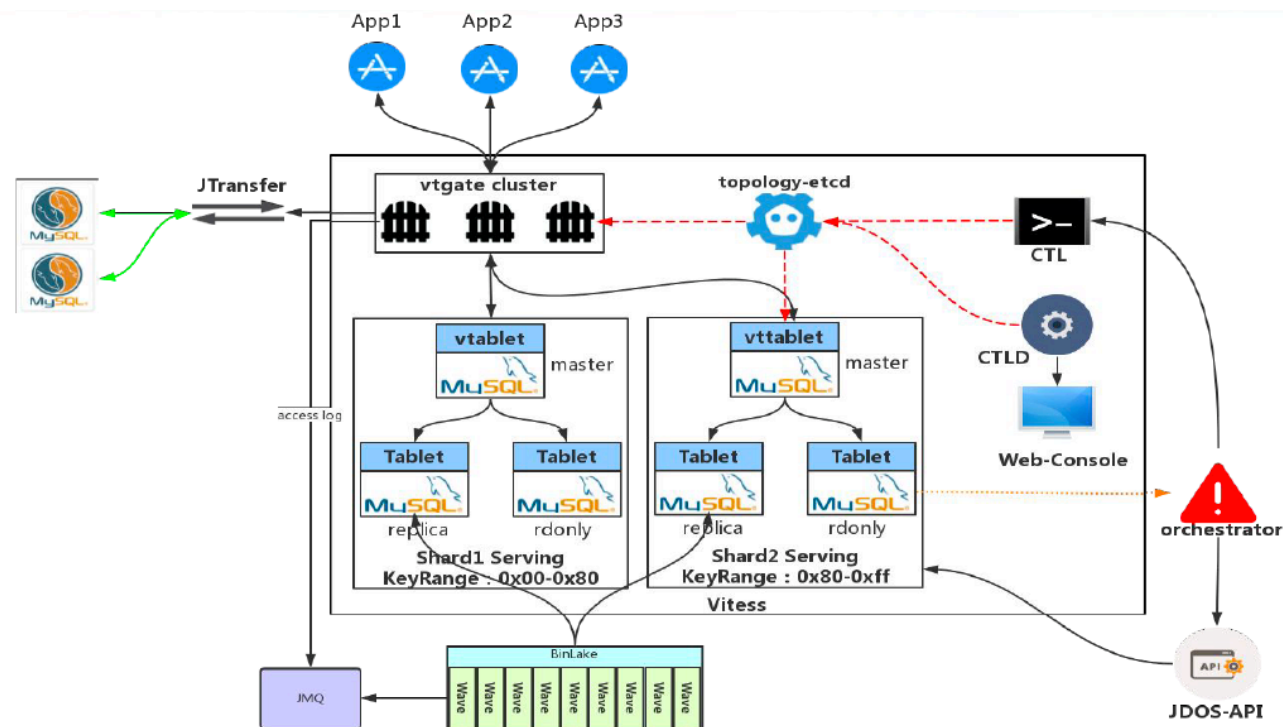
First Chinese internet company to make the Fortune Global 500

Strategic partnerships



JD Elastic Database

- MySQL for key businesses by serving large volumes of complex transactional data at JD.
- MySQL databases in containers on JD Kubernetes platform
 - ❑ Thousands of physical servers and tens of thousands of containers
- Vitess for scalable management and scaling.



One of the largest Vitess deployments

Problems and Challenges

- Difficult to estimate the resource demand
- Dynamic demand
- Performance guarantee
- Multiple dimensional resources: CPU, Memory, Disk, I/O Bandwidth
- Stateful applications with a lot of data

Workload Characterization

7 database systems and 631 containers


Application Name	Containers#	Days# (minutely)
DongDong	4	51
DistributedLog	6	51
OrderCancel	3	51
OrderTrace	49	42
TrainTicket	61	51
MallBill	192	51
SKUPrice	316	51

Table 1: Application details.

Key observations:

1. The 90-percentile utilization is about half or less of the maximum utilization.
2. Sizing based on 90-percentile utilization instead of the maximum utilization could lead to significant savings.

CPU Utilization



Right Sizing

- Estimate the workload demand

New workload

- Resource requirements
- Meta-data based classification + group resource requirements

Existing workload

- Historical usage analysis + prediction (ARIMA, LSTM, ...)

- Sizing based on percentile and correlation between workloads

- Use shared headroom to cope with peak demands
- Adjust both request and limit


Right Sizing of CPU Resources

Container

- Tail bound based guaranteed resource: $request = x\%ile$
- Maximum resources : $limit = y\%ile$
- Headroom: $headroom = limit - request$

Host

- Request: $Request = \sum_{i=1}^N request_i$
- Shared headroom
 - $Headroom \leq \sum_{i=1}^N headroom_i$
- Resource constraint:
 $Request + Headroom \leq Capacity$



Right Sizing of CPU Resources

Request : $Request = \sum_{i=1}^N request_i$

Shared headroom


- Non-correlated containers:

$$Headroom = MAX_i (limit_i - request_i)$$

- Correlated containers:

$$Headroom = SUM_i (limit_i - request_i)$$

Constraints: $Request + Headroom \leq Capacity$




$$Request = C1 + C2 + \dots + C9 \quad Headroom = MAX (H1 + H2 + H3, H4 + H5 + H6 + H7, H8 + H9)$$

$$Total = C1 + C2 + \dots + C9 + H4 + H5 + H6 + H7$$

$$max_of_sum < sum_of_max$$



Correlation Analysis



A large number of containers do not have strong correlation!

Experimental Evaluation


Comparison of Different Sizing Methods



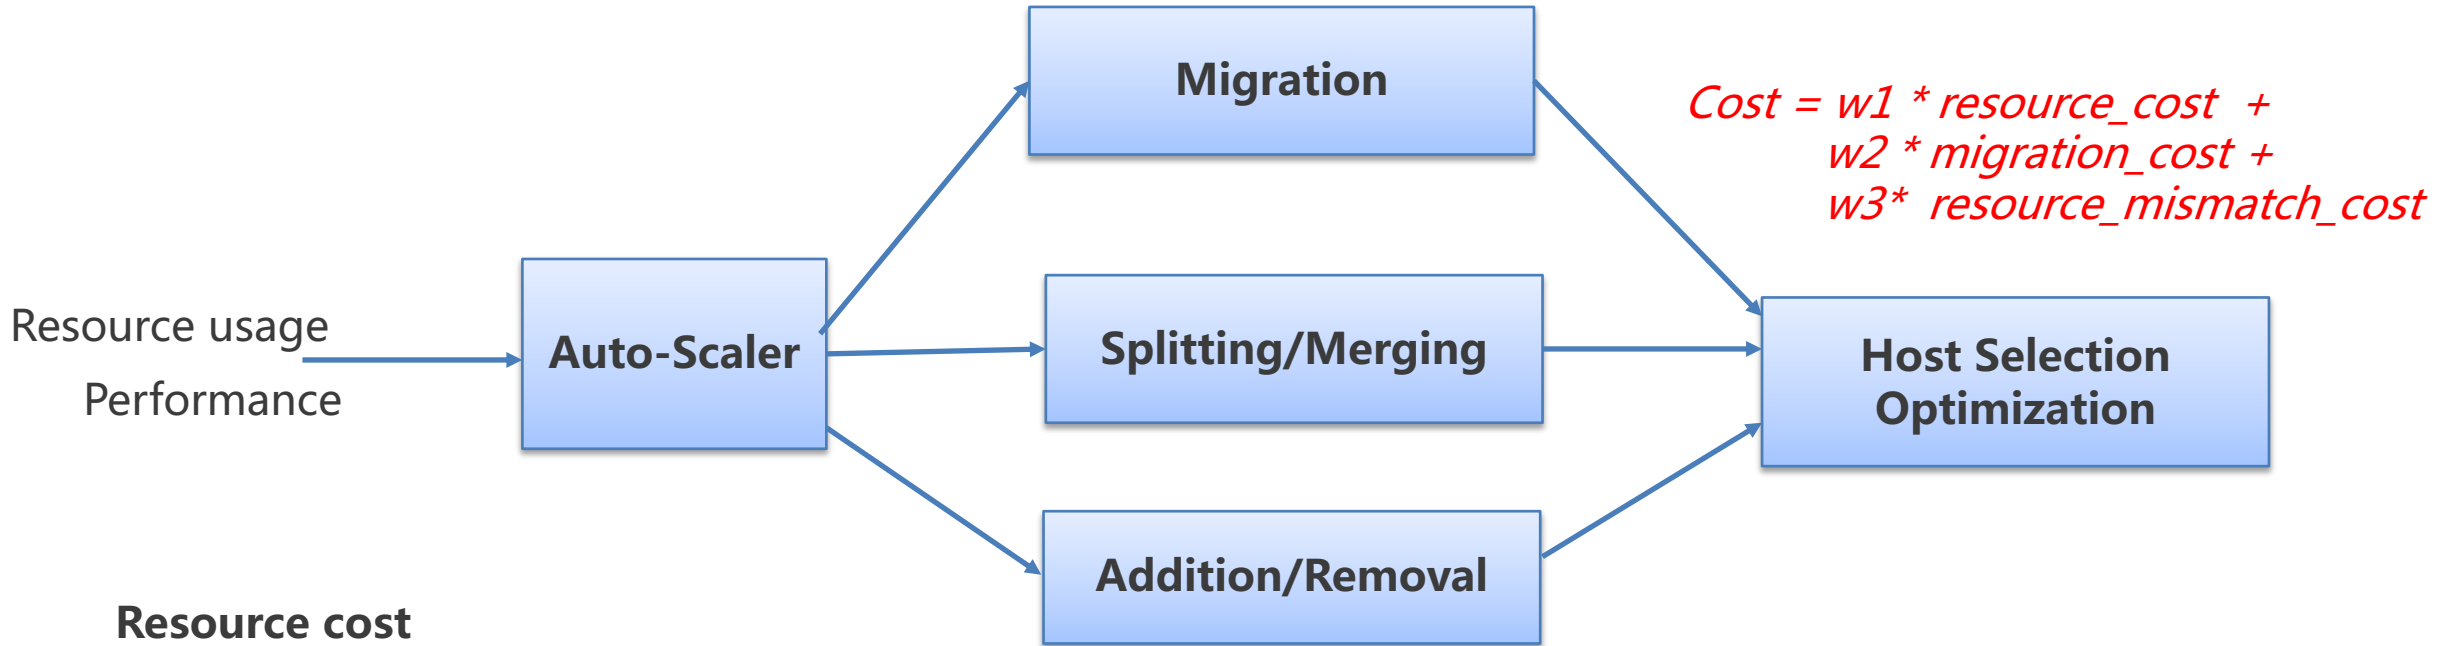
Randomly schedule 18 containers among 563 containers

Auto Scaling

- Split or merge instances
- Add or remove instances
- Migrate or consolidate instances



Auto-Scaling: Overview and Cost Models



Resource cost

$$\sum_{i=1}^M C_i (90\%ile) + \text{Max}_i \sum_{j=1}^{N_i} [C_{i,j} (99\%ile) - C_{i,j} (90\%ile)]$$

Migration cost: $k * \text{database size}$

Resource usability: *balancing, availability, anti-affinity*

Host Selection: Multi-Resource Balance

K8S: `BalancedResourceAllocation = 10 - abs(cpuRemainingFraction-memoryRemainingFraction)*10`


Optimization : similarity between the request resource and available resources

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}},$$

Host available resources : $(U_{\text{CPU}}, U_{\text{MEM}})$

Pod resource request : $(R_{\text{CPU}}, R_{\text{MEM}})$

Metric:
$$\frac{(R_{\text{CPU}} U_{\text{CPU}} + R_{\text{MEM}} U_{\text{MEM}})}{\sqrt{R_{\text{CPU}}^2 + R_{\text{MEM}}^2} \sqrt{U_{\text{CPU}}^2 + U_{\text{MEM}}^2}}$$




Experimental Evaluation

Setup

- 4,857 servers (8 configurations)
- 25,000 containers (120 configurations)

Results

- **Schedule 5%-20% more containers than K8S**
- **Comparable performance with K8S**



Host Selection: Resource Availability

K8S :

LeastRequestedPriority = $\text{cpu}((\text{capacity} - \text{sum}(\text{requested})) 10 / \text{capacity}) + \text{memory}((\text{capacity} - \text{sum}(\text{requested})) 10 / \text{capacity}) / 2$

Optimization: dynamic variable weights

$$w_{\text{cpu}} * \text{cpuAvailFraction} + w_{\text{mem}} * \text{memAvailFraction}$$

Algorithm 1:

Weighted sum of CPU : $\text{cpuFraction} = \sum_i f_i * \text{cpuFraction}_i$

Weight sum of memory : $\text{memFraction} = \sum_i f_i * \text{memFraction}_i$

$$w_{\text{cpu}} = \frac{\text{cpuFraction}}{\text{cpuFraction} + \text{memFraction}}$$
$$w_{\text{mem}} = \frac{\text{memFraction}}{\text{cpuFraction} + \text{memFraction}}$$

Algorithm 2: Hosts with stable resource usage are more usable.

MAD- Median Absolute Deviation

$$w_{\text{cpu}} = 1 - \text{median}(|\text{cpuUtil}_i - \text{median}(\text{cpuUtil}_i)|)$$

$$\cdot w_{\text{mem}} = 1 - \text{median}(|\text{memUtil}_i - \text{median}(\text{memUtil}_i)|)$$

Host Selection: Correlation-awareness

X : existing pods resource usage

Y : new pod resource usage

Anti-affinity: If Y can be predicted by Y, X and Y are highly correlated. We should avoid placing them on the same host.

$$\mathbf{X} = \begin{bmatrix} 1 & x_{1,1} & \cdots & x_{1,k} & \cdots & x_{1,K} \\ \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_{n,1} & \cdots & x_{n,k} & \cdots & x_{n,K} \\ \vdots & \vdots & \ddots & \vdots & \ddots & \vdots \\ 1 & x_{N,1} & \cdots & x_{N,k} & \cdots & x_{N,K} \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_n \\ \vdots \\ y_N \end{bmatrix}$$

Multi-regression analysis $\mathbf{b} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad \hat{\mathbf{y}} = \mathbf{X} \mathbf{b}$

Correlation
$$\frac{\sum_{i=1}^n (y_i - m_Y)^2 (\hat{y}_i - m_{\hat{Y}})^2}{\sum_{i=1}^n (y_i - m_Y)^2 \sum_{i=1}^n (\hat{y}_i - m_{\hat{Y}})^2}$$

Conclusions

- Running MySQL in containers on Kubernetes and using Vitess cluster management is the foundation of MySQL resource optimization and management.
- Statistical analysis and prediction based resource sizing and scaling approaches are useful to improve resource utilization of containerized MySQL.
- **Kubernetes + Vitess + Advanced Resource Optimization Algorithms** have significantly improved the resource efficiency and reduced the operations and maintenance costs of running large scale MySQL clusters at JD.



Acknowledgements

Thanks for the following colleagues!

Haifeng Liu, Jiangang Fan, Fengcai Liu, Guangya Bao, Huasong Shan, An Peng, Liu Yan

Thank you for your time!

Contact:

Yuan Chen

Email: yuan.chen@jd.com

WeChat: yuan_gt



System Architecture

JD.COM 京东

